

Informed Search

Lecture 3

Uninformed and informed searches

- Based on the search problems we can classify the search algorithms into two major types:
 - Uninformed search (Blind Search) Algorithms.
NB: “covered in the previous lecture”
 - Informed search (Heuristic Search) Algorithms.
NB: “will be covered shortly”

Uninformed Search algorithms

- These kind of algorithms does not contain any Knowledge of the domain such as how closeness to the goal or the location of the goal.
- They operate in a brute force way as it only includes information about how to traverse the tree and how to identify leaf and goals.
- Applies away in which search tree is searched without any information about the goal.
- They examine each node until it reaches the goal node and it can stop after that.

Examples

- Depth first search (DFS)
- Breadth first search (BFS)
- Uniform cost search (UCS)
- Depth limited search
- Iterative Deepening depth first search
- Bidirectional search

Informed search Algorithms.

- They use domain knowledge, the problem information to the goal is available which can guide the search.
- Informed search strategies can find a solution more efficient than an uninformed search.
- Heuristic is away which might not always be guaranteed for the best solution but guaranteed to find a good solution in a reasonable time.
- They can solve much complex problem which could not be solved in other ways.

Examples

- Greedy Best First Search
- A* Search
- Hill climbing Search

Heuristic Function

- Is a function that estimates how close is a state to a goal
- Designed for a particular search problem
i.e: May differ from one problem to another
- Example: Euclidean distance or Manhattan distance for a path problem
- The accuracy of choosing the heuristic function affects the accuracy of the algorithm.

A heuristic function

- Let evaluation function $h(n)$ (**h**euristic)
 - $h(n) = \textit{estimated}$ cost of the cheapest path from node n to goal node.
 - If n is goal then $h(n)=0$

Best First Search

- Uses an evaluation function $f(n)$ for each node and the node to be expanded is the node n with the smallest $f(n)$
- Example
 - A* Search
 - Greedy best first search

A Quick Review

- $g(n)$ = cost from the initial state to the current state n
- $h(n)$ = estimated cost of the cheapest path from node n to a goal node
- $f(n)$ = evaluation function to select a node for expansion (usually the lowest cost node)

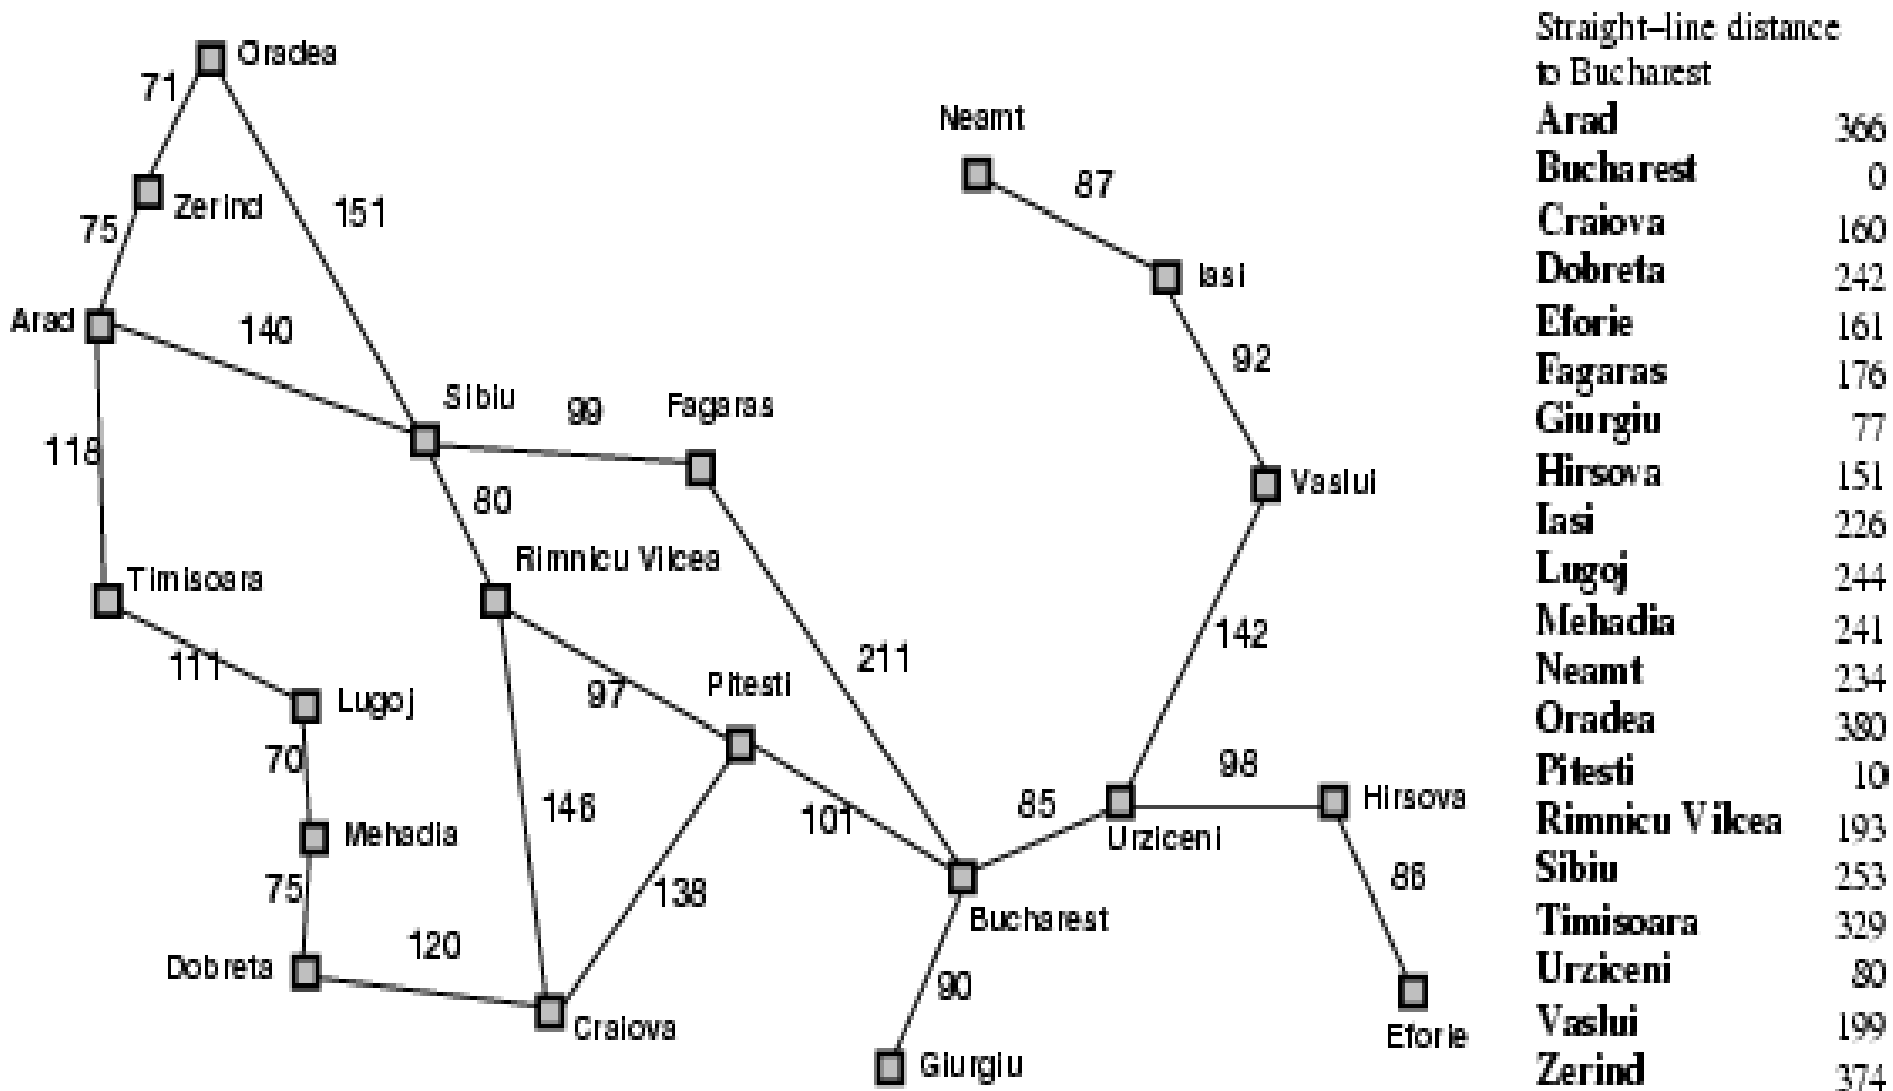
Examples

- Greedy Best First Search
- A* Search
- Hill climbing Search


Greedy best-first search

- Evaluation function $f(n) = h(n)$
(**h**euristic)
= estimate of cost from n to *goal*
- Ignores the path cost
- Greedy best-first search expands the node that **appears** to be closest to goal


Romania with step costs in km





Greedy best-first search example




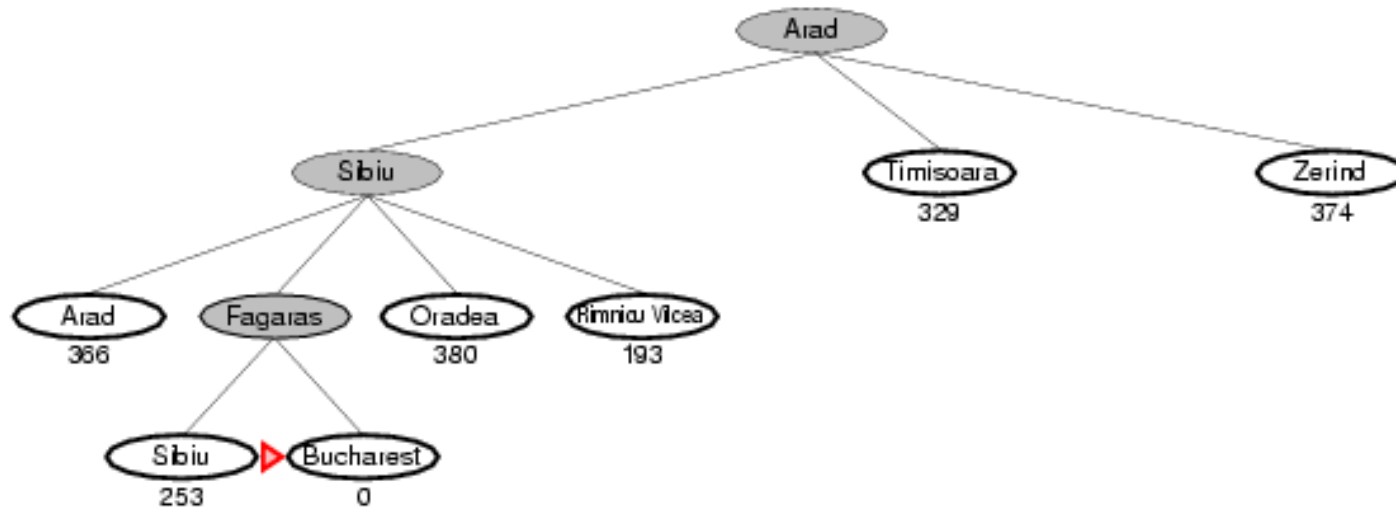
Greedy best-first search example




Greedy best-first search example



Greedy best-first search example



Optimal Path



Greedy Best-First Search Algorithm

Input: State Space

Output: *failure* or path from a start state to a goal state.


Assumptions:

- L is a list of nodes that have not yet been examined ordered by their h value.
- The state space is a tree where each node has a single parent.


1. Set L to be a list of the initial nodes in the problem.
 2. While L is not empty
 1. Pick a node n from **the front of** L .
 2. If n is a goal node
 1. stop and return it and the path from the initial node to n .
 - Else
 1. remove n from L .
 2. For each child c of n
 1. insert c into L **while preserving the ordering of nodes in** L and labelling c with its path from the initial node as well as its h value.
 - End for
 - End if
 - End while
- Return *failure*

Greedy BF Search

- Strategy: expand a node that you think is closest to a goal state
 - Heuristic: estimate of distance to nearest goal for each state



- A common case:
 - Best-first takes you straight to the (wrong) goal



- Worst-case: like a badly-guided DFS

Properties of greedy best-first search

- Complete?
 - Not unless it keeps track of all states visited
 - Otherwise can get stuck in loops (just like DFS)
- Optimal?
 - No – we just saw a counter-example

Examples

- Greedy Best First Search
- **A* Search**
- Hill climbing Search

A* Search Algorithm

Evaluation function $f(n) = h(n) + g(n)$

$h(n)$ estimated cost to goal from n

$g(n)$ cost so far to reach n

A* uses admissible heuristics, i.e.,

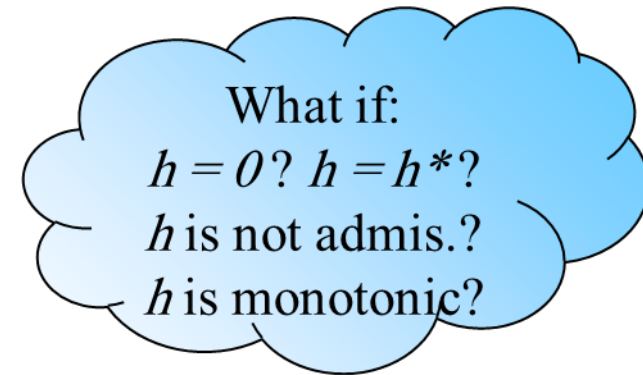
$h(n) \leq h^*(n)$ where $h^*(n)$ is the

true cost from n .

A* Search finds the
optimal path

A* search

- Best-known form of best-first search.
- Idea: avoid expanding paths that are already expensive.
- Combines uniform-cost and greedy search
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ the cost (so far) to *reach* the node
 - $h(n)$ estimated cost to *get from the node to the goal*
 - $f(n)$ estimated *total cost* of path through n to goal
- Implementation: Expand the node n with minimum $f(n)$




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 1, expanded: 1

expnd. node	Frontier
	{S:8}
S not goal	{A:1+8,B:5+4,C:8+3}




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 2, expanded: 2

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A not goal	{B:9,G:1+9+0,C:11, D:1+3+ ∞ ,E:1+7+ ∞ }




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 3, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D:∞,E:∞}
B not goal	{G:5+4+0, C:10 , C:11, D:∞,E:∞} replace




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 4, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D:∞,E:∞}
B	{G:9,C:11,D:∞,E:∞}
G goal	{C:11,D:∞,E:∞} not expanded




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 4, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D:∞,E:∞}
B	{G:9,C:11,D:∞,E:∞}
G	{C:11,D:∞,E:∞}




- Pretty fast and optimal*

path: S,B,G
cost: 9


A* search example



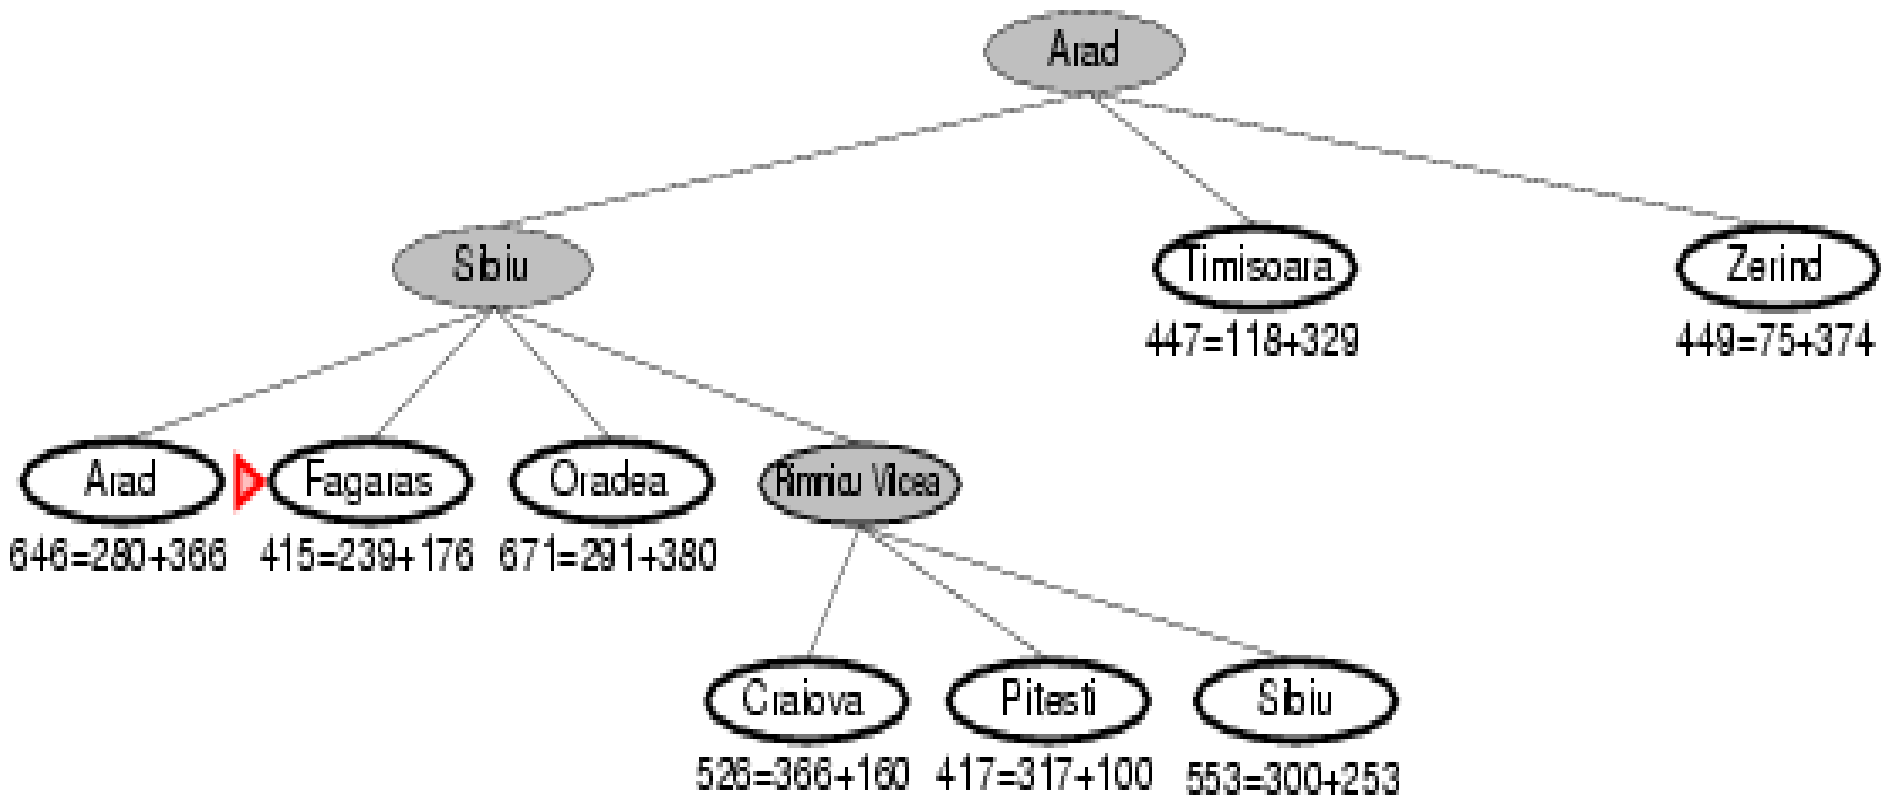
A* search example



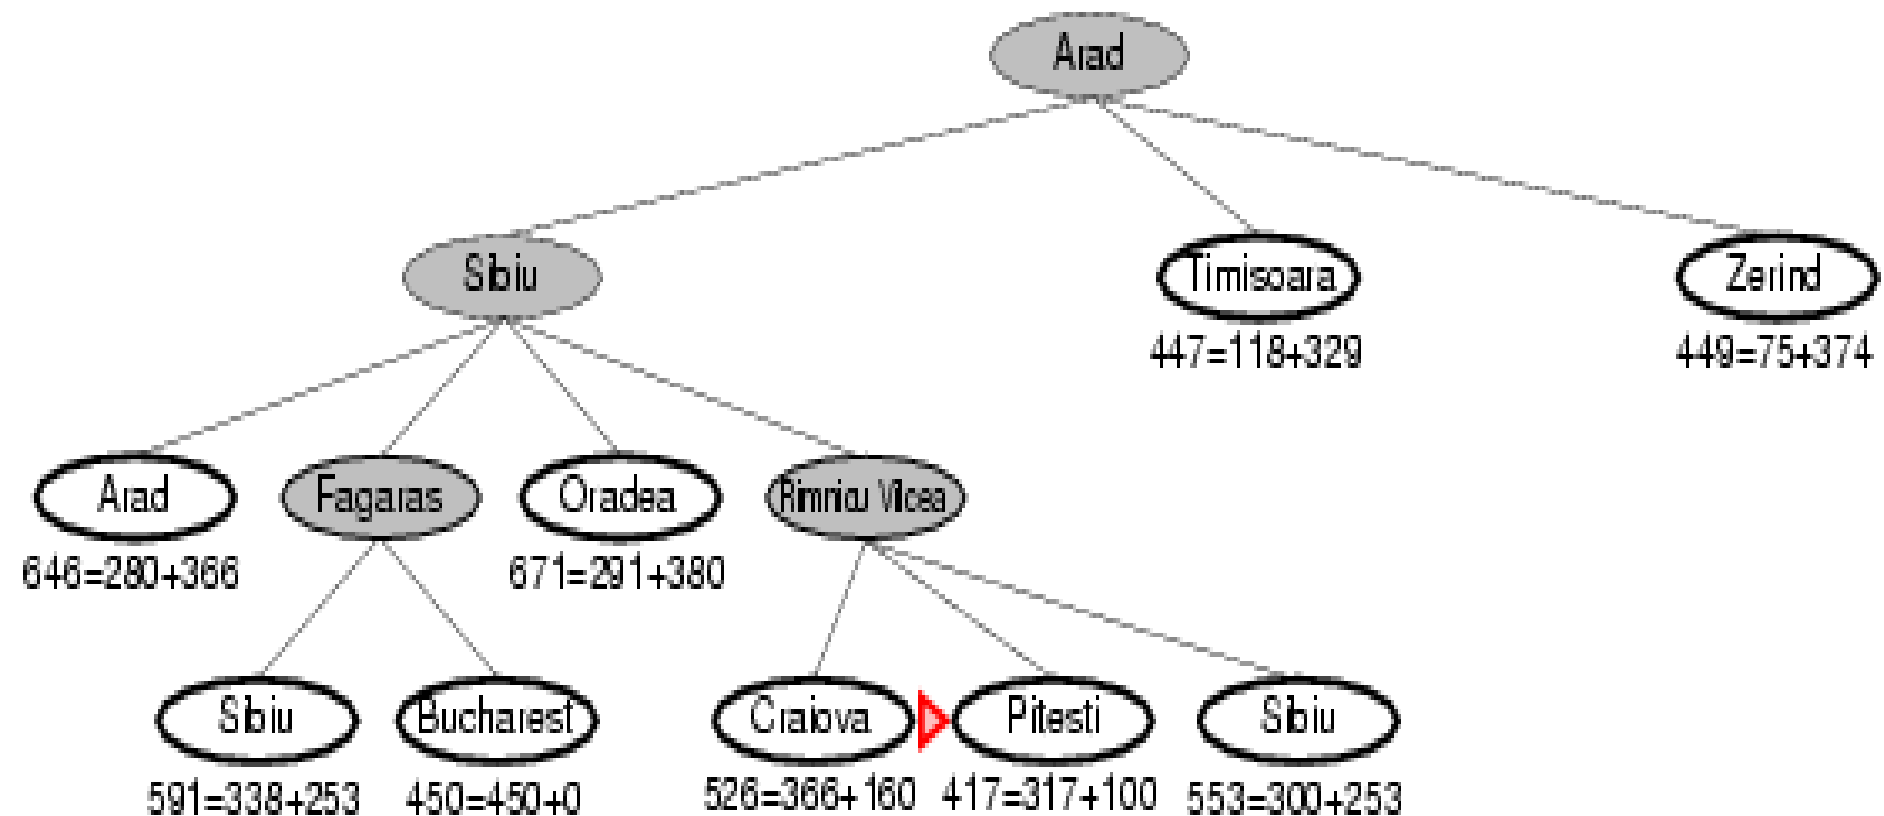
A* search example



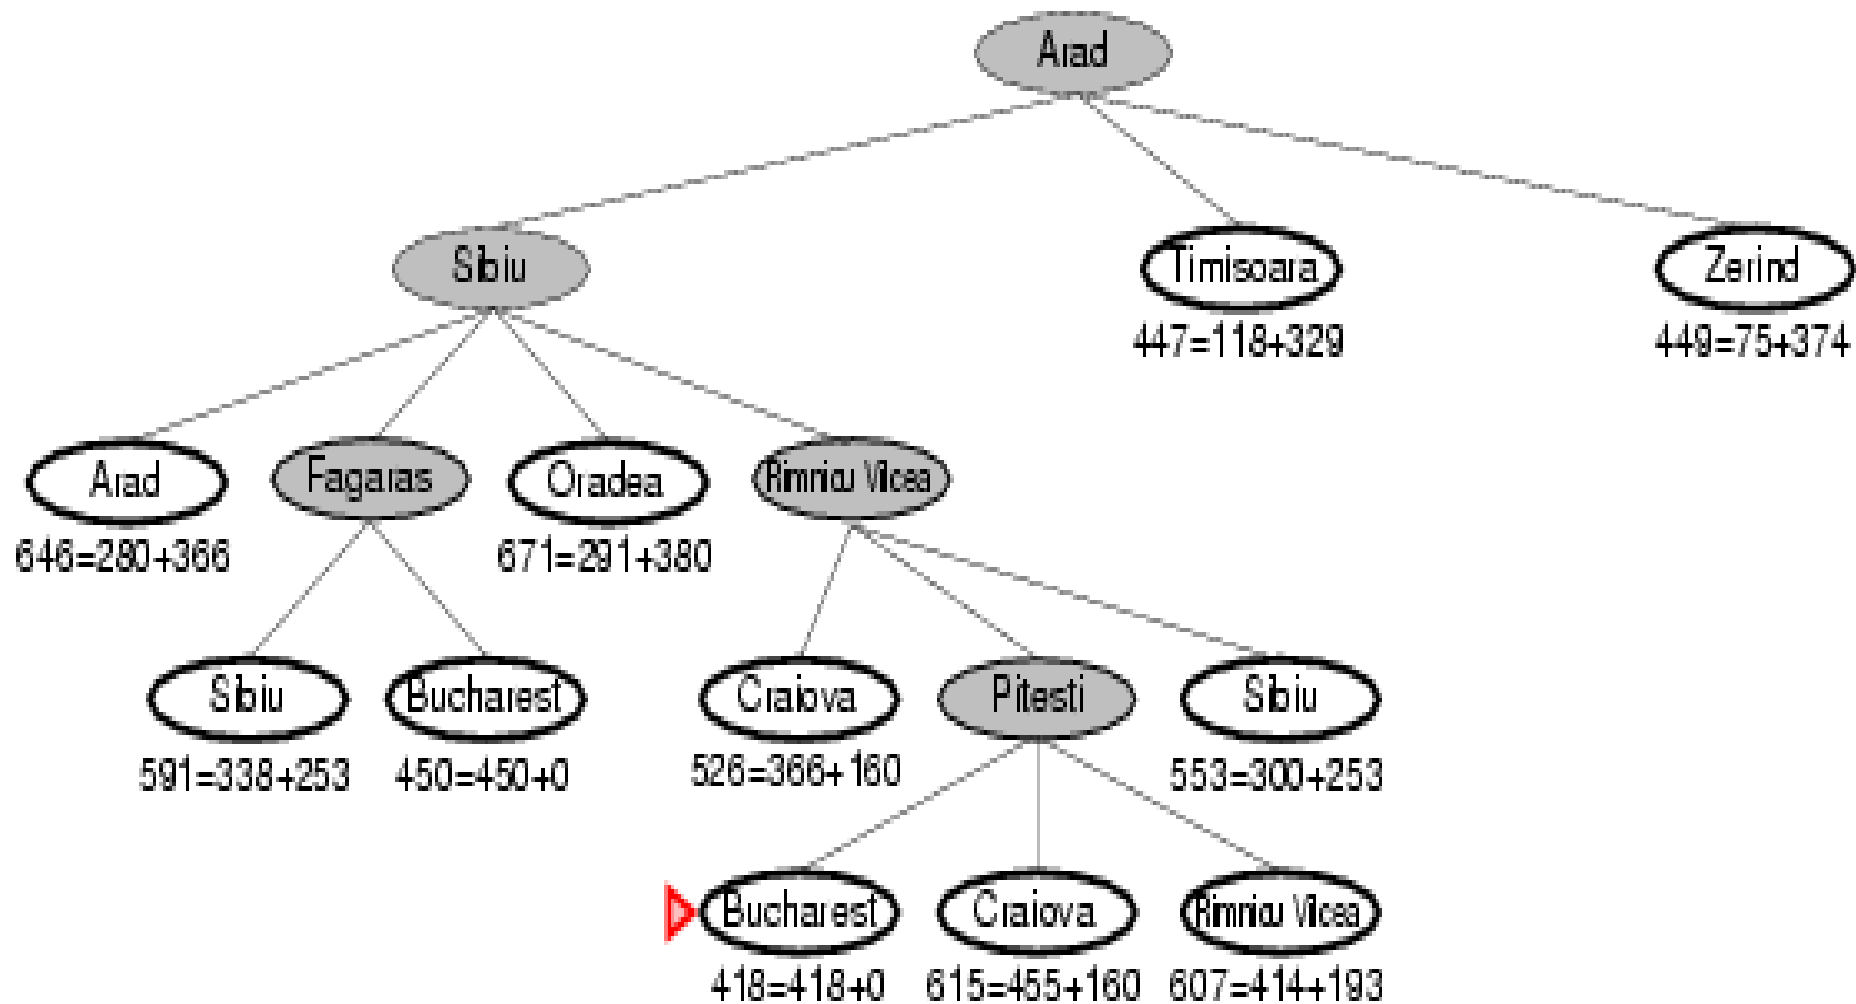
A* search example




A* search example



A* search example



$f(n) = g(n) + h(n)$
 $g(n) = \text{distance from start}$
 $h(n) = \text{\# tiles out of place}$



Properties of A^*

- Complete? Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- Optimal? Yes

A*: Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



Examples


- Greedy Best First Search
- A* Search
- Hill climbing Search

Hill Climbing Search

- For artefact-only problems (don't care about the path)
- Depends on some $e(\text{state})$
 - Hill climbing tries to maximise score e
- Randomly choose a state
 - Only choose actions which improve e
 - If cannot improve e , then perform a **random restart**
 - Choose another random state to restart the search from
- Only ever have to store one state (the present one)
 - Can't have cycles as e always improves

Hill-climbing search


- Problem: depending on initial state, can get stuck in local maxima



Hill Climbing - Algorithm

1. Pick a random point in the search space
2. Consider all the neighbors of the current state
3. Choose the neighbor with the best quality and move to that state
4. Repeat 2 thru 4 until all the neighboring states are of lower quality
5. Return the current state as the solution state

Example: 8 Queens



- Place 8 queens on board
 - So no one can “take” another
- Gradient descent search
 - Throw queens on randomly
 - e = number of pairs which can attack each other
 - Move a queen out of other’s way
 - Decrease the evaluation function
 - If this can’t be done
 - Throw queens on randomly again

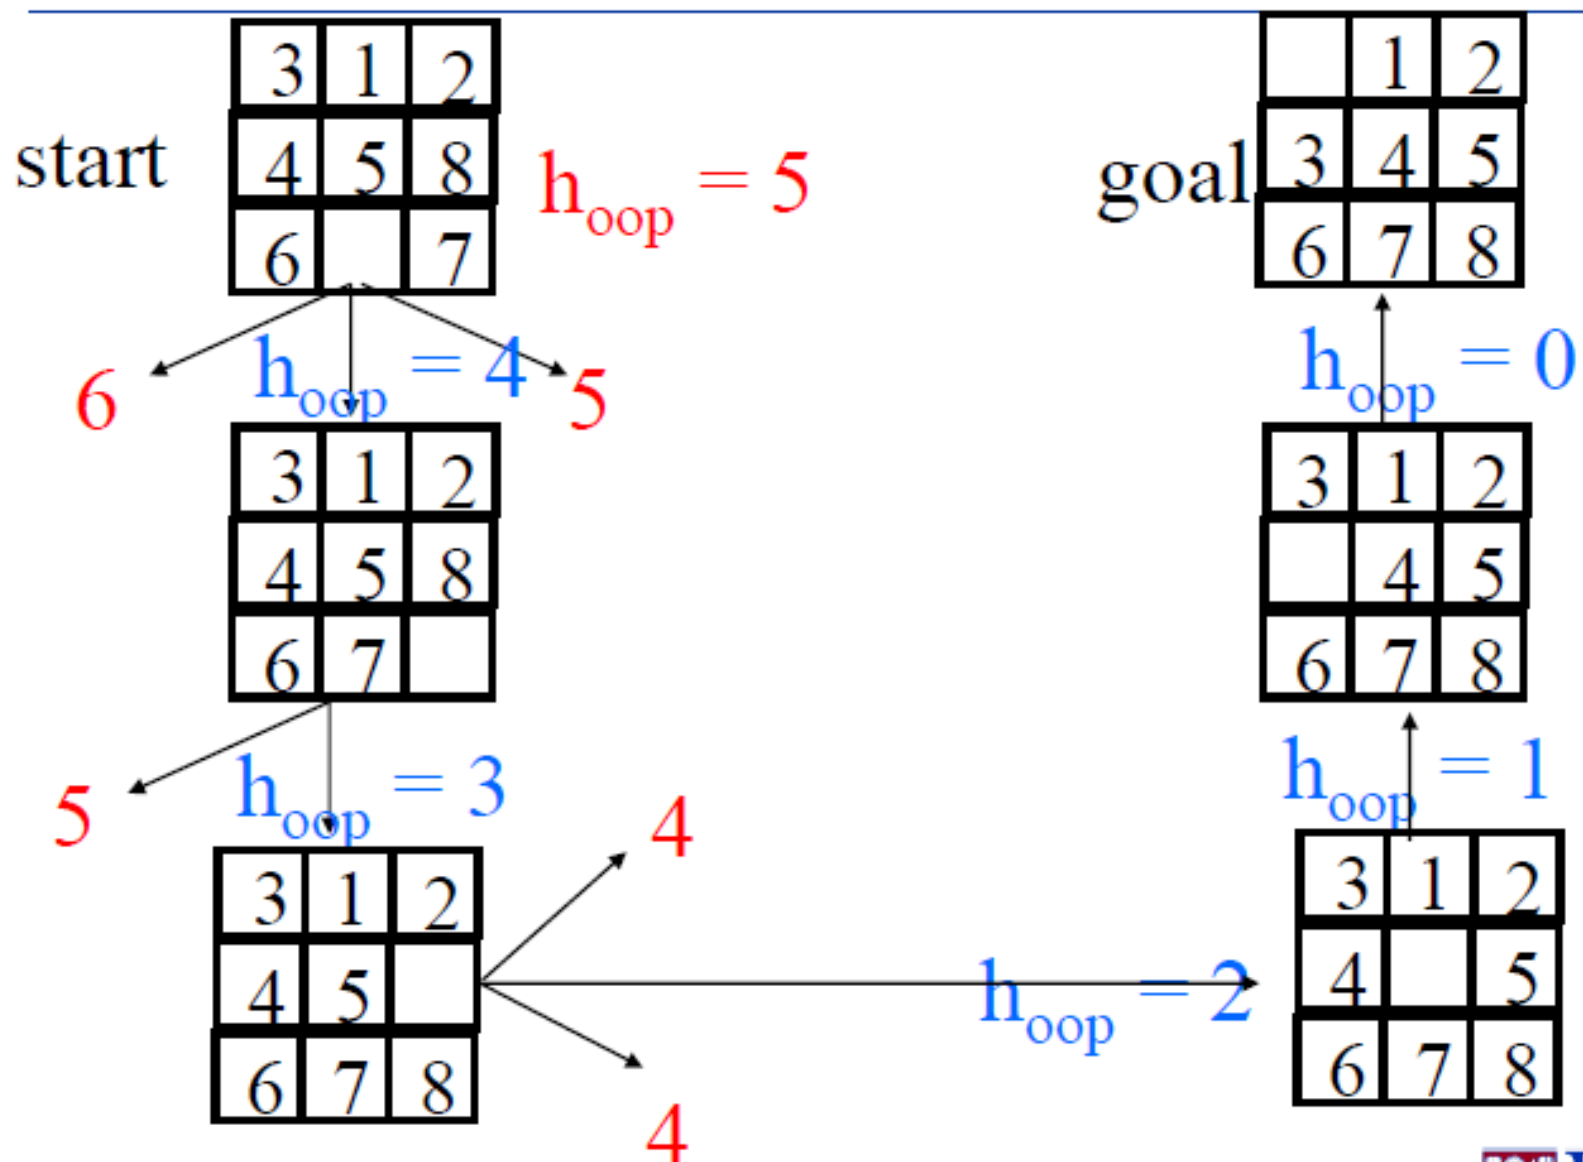
Hill-climbing search

- Looks one step ahead to determine if any successor is better than the current state; if there is, move to the best successor.
- **Rule:** If there exists a successor s for the current state n such that
 - $h(s) < h(n)$ and
 - $h(s) \leq h(t)$ for all the successors t of n ,then move from n to s . Otherwise, halt at n .

Hill-climbing search

- Similar to Greedy search in that it uses $h()$, but does not allow backtracking or jumping to an alternative path since it doesn't "remember" where it has been.
- If Hill climbing failed to find the goal it will start from another random node searching for the goal node(may choose the same failed node ☹️).

Hill climbing example I (*minimizing h*)



Properties of Hill Climbing

- Complete? NO
- Optimal? NO